

# Machine Learning

CS 539

Worcester Polytechnic Institute

Department of Computer Science

Instructor: Prof. Kyumin Lee

# Midterm Review

# Midterm Exam

- The exam will be held at 4pm in the next class, Feb 27.
- Bring your laptop/phone so that you can submit your answers via Canvas.
- Remote exam will not be allowed.
- The exam is closed book.
- You may prepare and use one standard 8.5" by 11" piece of paper with any notes you think appropriate or significant (use only \*one-side\*).
- You may use a calculator if it make you feel comfortable. But no other electronic devices are allowed (e.g., cell phone, tablet and computer). You will use your laptop/phone for answer submission only!

# Previous Class...

Supervised Learning:

→ Regression and  
Classification

Decision Tree

→ How to choose the  
best attribute?  
→ Information Gain  
with Entropy

# Previous Class...

Overfitting in Decision Tree? How to avoid it?

→ Reduced error pruning,

kNN

→ 1-NN vs k-NN

→ How to reduce the prediction time?

Evaluation

→ Accuracy, Precision and Recall

# Previous Class...

Cross-validation

Linear Regression  
and  
Regularization

Perceptron

# Previous Class...

Logistic Regression

VC dimension  
Bias vs Variance

Neural Network  
→ Feedforward  
propagation  
→ Back propagation

# Upcoming Schedule

- Midterm
  - Feb 27 in class
- HW4 will be out on March 1st (postponed)
- Project Workday on March 12 and 15
- Project Proposal
  - Due date is March 18



# Project Workday


Date	Presenters
3/12	<ul style="list-style-type: none"><li>• Yu-Chi Liang, William Ryan, Riley Blair, and Stephen Fanning</li><li>• Chris Lee, Andrew Kerekon, Amulya Mohan, Alex Siracusa, and Sulaiman Moukheiber</li><li>• Vagmi Bhagavathula, Deepti Gosukonda, Adina Palayoor, Bishoy Soliman Hanna, and Jared Chan</li><li>• Sreeram Marimuthu, Oruganty Nitya Phani Santosh, Sarah Olson, and Thomas Pianka</li><li>• Shubham Dashrath Wagh, Atharva Pradip Kulkarni, Amit Virchandbhai Prajapati, Niveditha Narasimha Murthy</li><li>• Aria Yan, Alisha Peeriz, Nupur Kalkumbe, Pavan Antala, Rutuja Madhumilind Dongre</li><li>• Noushin Khosravi Largani, Jinqin Xiong, Kexin Li, Ronit Kapoor, and Yiqun Duan</li><li>• Phil Brush, Liam Hall, Jared Morgan, Alex</li></ul>
3/15	<ul style="list-style-type: none"><li>• Khang Luu, Austin Aguirre, Brock Dubey, Ivan Klevanski,</li><li>• Adhiraj, Karl, Shariq Madha, Yue Bao, Vasilli Gorbunov</li><li>• Edward Smith, Michael Alicea, Cutter Beck, Blake Bruell, Anushka Bangal</li><li>• Rahul Chhatbar, Sonu Tejwani, Deep Suchak, Shoan Bhatambare</li><li>• Daniel Fox, Bijesh Shrestha, Chad Hucey, Aayush Sangani, Ivan Lim</li><li>• Devesh Bhangale, Shipra Poojary, Jagruti Chitte, Parth Shroff, Saurabh Pande</li><li>• Alessandra Serpes, Khushita Joshi, Sanjeeeth Nagappa Chakrasali, Shaun Noronha, Sankalp Vyas</li><li>• Rohan Rana, Theo Coppola, Olivia Raisbeck</li></ul>

More detailed logistics will be announced on March 1st

# Convolutional Neural Network (CNN)

## **Convolution:**

Convolution is a mathematical operation on two functions to produce a third function that expresses how the shape of one is modified by the other. It is a term that the neural network community (Hinton, specifically) adopted from the signal and image processing communities to architect the “feature detection” layers of “deep” neural networks.



1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



1	0	0	0	0
0	1	0	0	1
0	0	1	1	0
1	0	0	0	1
0	1	0	0	1
0	0	1	0	1



6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1


1: 1  
 2: 0  
 3: 0  
 4: 0  
 ⋮  
 7: 0  
 8: 1  
 9: 0  
 10: 0  
 ⋮  
 13: 0  
 14: 0  
 15: 1  
 16: 1  
 ⋮

3











Less parameters!






6 x 6 image

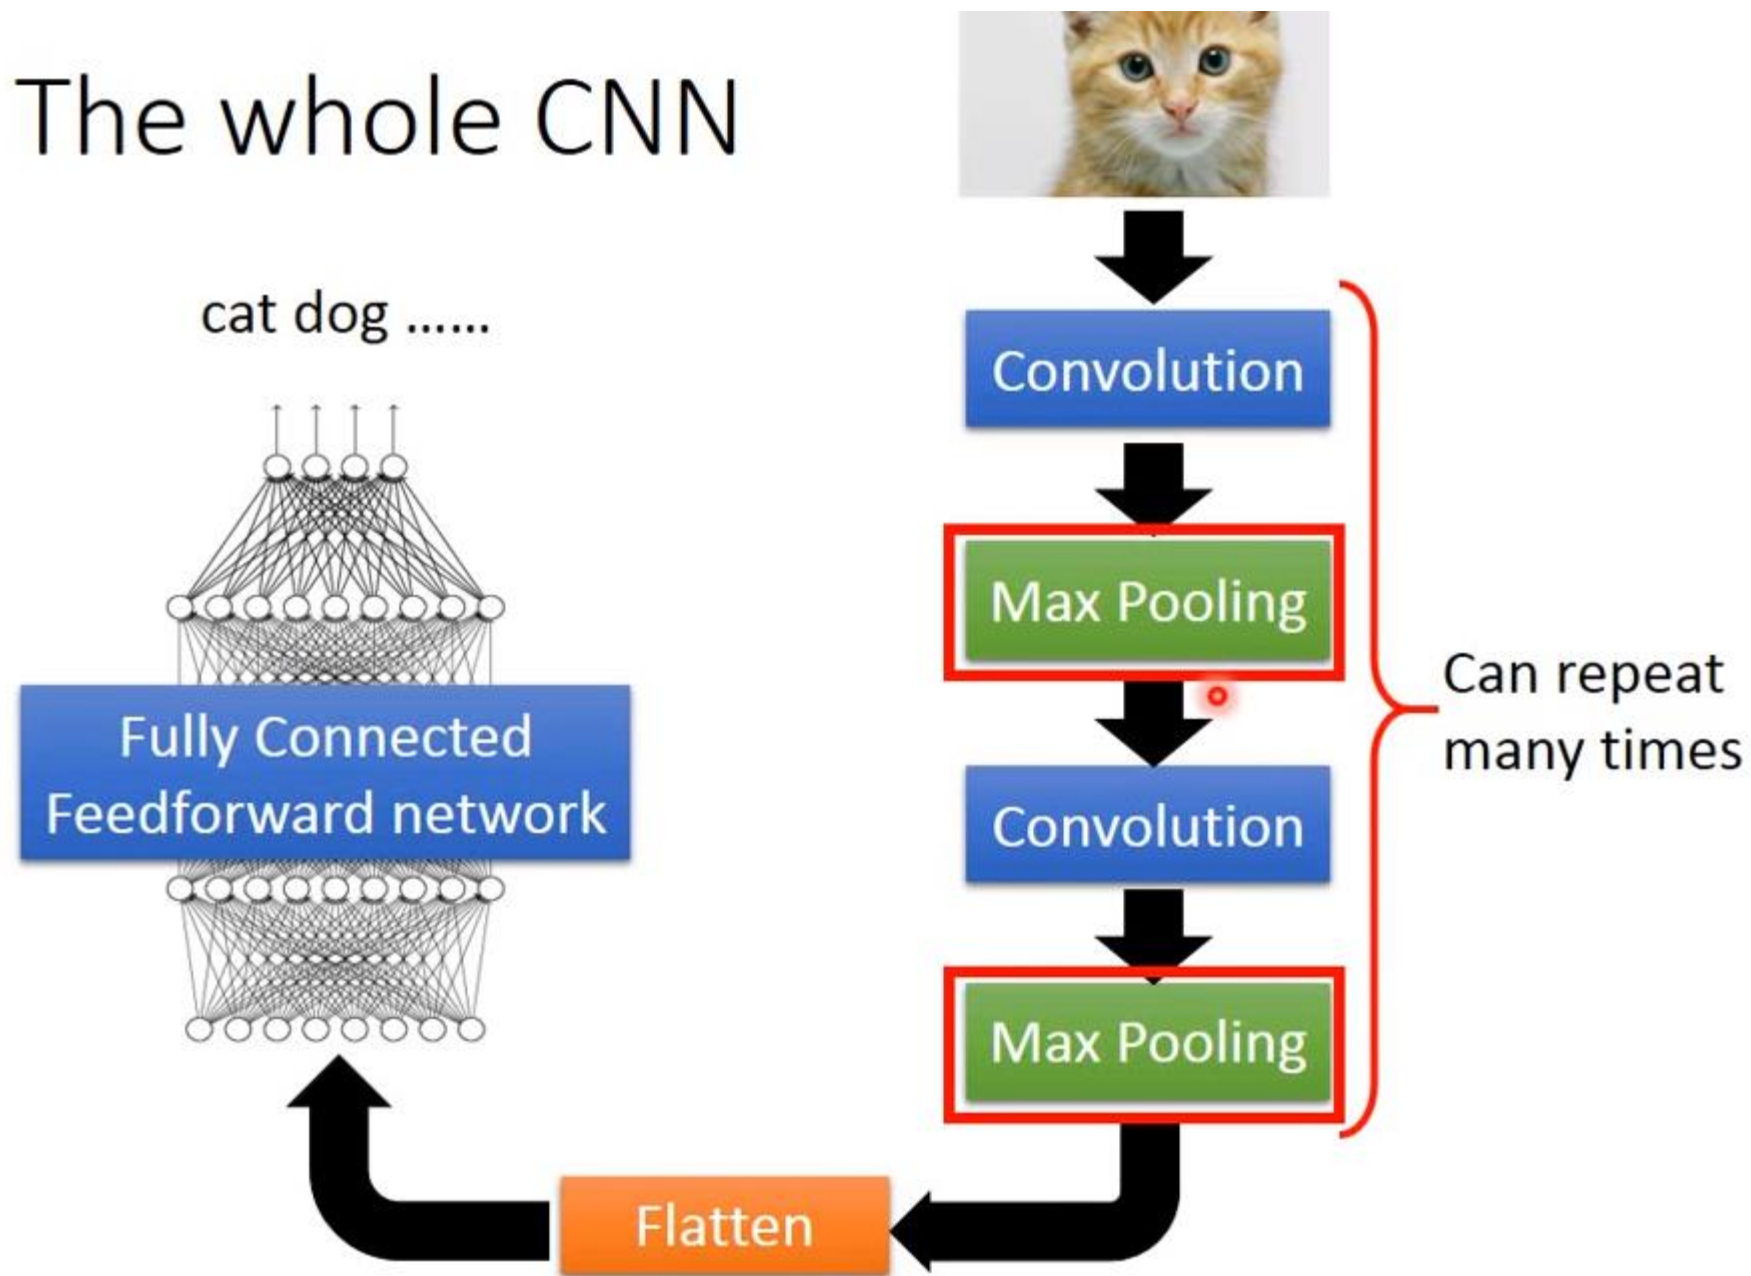
Less parameters!

Even less parameters!





# The whole CNN



# VGGNet

Small filters, Deeper networks

16 - 19 layers (VGG16Net)

Only 3x3 CONV stride 1, pad 1  
and 2x2 MAX POOL stride 2



# VGGNet

INPUT: [224x224x3]    memory:  $224*224*3=150K$     params: 0    (not counting biases)

CONV3-64: [224x224x64]    memory:  $224*224*64=3.2M$     params:  $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64]    memory:  $224*224*64=3.2M$     params:  $(3*3*64)*64 = 36,864$

POOL2: [112x112x64]    memory:  $112*112*64=800K$     params: 0

CONV3-128: [112x112x128]    memory:  $112*112*128=1.6M$     params:  $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128]    memory:  $112*112*128=1.6M$     params:  $(3*3*128)*128 = 147,456$

POOL2: [56x56x128]    memory:  $56*56*128=400K$     params: 0

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*256)*256 = 589,824$

POOL2: [28x28x256]    memory:  $28*28*256=200K$     params: 0

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512]    memory:  $14*14*512=100K$     params: 0

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$


CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512]    memory:  $7*7*512=25K$     params: 0

FC: [1x1x4096]    memory: 4096    params:  $7*7*512*4096 = 102,760,448$

FC: [1x1x4096]    memory: 4096    params:  $4096*4096 = 16,777,216$

FC: [1x1x1000]    memory: 1000    params:  $4096*1000 = 4,096,000$



VGG16

# VGGNet

INPUT: [224x224x3]    memory:  $224*224*3=150K$     params: 0    (not counting biases)

CONV3-64: [224x224x64]    memory:  $224*224*64=3.2M$     params:  $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64]    memory:  $224*224*64=3.2M$     params:  $(3*3*64)*64 = 36,864$

POOL2: [112x112x64]    memory:  $112*112*64=800K$     params: 0

CONV3-128: [112x112x128]    memory:  $112*112*128=1.6M$     params:  $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128]    memory:  $112*112*128=1.6M$     params:  $(3*3*128)*128 = 147,456$

POOL2: [56x56x128]    memory:  $56*56*128=400K$     params: 0

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*256)*256 = 589,824$

POOL2: [28x28x256]    memory:  $28*28*256=200K$     params: 0

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512]    memory:  $14*14*512=100K$     params: 0

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512]    memory:  $7*7*512=25K$     params: 0


FC: [1x1x4096]    memory: 4096    params:  $7*7*512*4096 = 102,760,448$

FC: [1x1x4096]    memory: 4096    params:  $4096*4096 = 16,777,216$

FC: [1x1x1000]    memory: 1000    params:  $4096*1000 = 4,096,000$

**TOTAL memory:**  $24M * 4 \text{ bytes} \approx 96MB$  / image (for a forward pass)

**TOTAL params:** 138M parameters



VGG16

# VGGNet

INPUT: [224x224x3] memory:  $224*224*3=150K$  params: 0 (not counting biases)

CONV3-64: [224x224x64] memory:  $224*224*64=3.2M$  params:  $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory:  $224*224*64=3.2M$  params:  $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory:  $112*112*64=800K$  params: 0

CONV3-128: [112x112x128] memory:  $112*112*128=1.6M$  params:  $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory:  $112*112*128=1.6M$  params:  $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory:  $56*56*128=400K$  params: 0

CONV3-256: [56x56x256] memory:  $56*56*256=800K$  params:  $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory:  $56*56*256=800K$  params:  $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory:  $56*56*256=800K$  params:  $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory:  $28*28*256=200K$  params: 0

CONV3-512: [28x28x512] memory:  $28*28*512=400K$  params:  $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory:  $28*28*512=400K$  params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory:  $28*28*512=400K$  params:  $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory:  $14*14*512=100K$  params: 0

CONV3-512: [14x14x512] memory:  $14*14*512=100K$  params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory:  $14*14*512=100K$  params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory:  $14*14*512=100K$  params:  $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory:  $7*7*512=25K$  params: 0

FC: [1x1x4096] memory: 4096 params:  $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params:  $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params:  $4096*1000 = 4,096,000$

Note:

Most memory is in  
early CONV

Most params are  
in late FC

TOTAL memory:  $24M * 4 \text{ bytes} \sim 96MB$  / image (only forward!  $\sim 2$  for bwd)

TOTAL params: 138M parameters



# VGGNet

INPUT: [224x224x3]      memory:  $224*224*3=150K$     params: 0      (not counting biases)

CONV3-64: [224x224x64]    memory:  $224*224*64=3.2M$     params:  $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64]    memory:  $224*224*64=3.2M$     params:  $(3*3*64)*64 = 36,864$

POOL2: [112x112x64]    memory:  $112*112*64=800K$     params: 0

CONV3-128: [112x112x128]    memory:  $112*112*128=1.6M$     params:  $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128]    memory:  $112*112*128=1.6M$     params:  $(3*3*128)*128 = 147,456$

POOL2: [56x56x128]    memory:  $56*56*128=400K$     params: 0

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256]    memory:  $56*56*256=800K$     params:  $(3*3*256)*256 = 589,824$

POOL2: [28x28x256]    memory:  $28*28*256=200K$     params: 0

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512]    memory:  $28*28*512=400K$     params:  $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512]    memory:  $14*14*512=100K$     params: 0

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512]    memory:  $14*14*512=100K$     params:  $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512]    memory:  $7*7*512=25K$     params: 0

FC: [1x1x4096]    memory: 4096    params:  $7*7*512*4096 = 102,760,448$

FC: [1x1x4096]    memory: 4096    params:  $4096*4096 = 16,777,216$

FC: [1x1x1000]    memory: 1000    params:  $4096*1000 = 4,096,000$

**TOTAL memory: 24M \* 4 bytes ~= 96MB / image** (only forward! ~\*2 for bwd)

**TOTAL params: 138M parameters**




VGG16

Common names

[Zeiler and Fergus 2013]

Visualization of VGG-16 by Lane McIntosh. VGG-16 architecture from [Simonyan and Zisserman 2014].




Low-level features


Mid-level features

High-level features


Linearly separable classifier



VGG-16 Conv1\_1



VGG-16 Conv3\_2




VGG-16 Conv5\_3



# Implementation of CNN in PyTorch

For this tutorial, we will use the CIFAR10 dataset. It has the classes: 'airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck'. The images in CIFAR-10 are of size 3x32x32, i.e. 3-channel color images of 32x32 pixels in size.



cifar10

[https://pytorch.org/tutorials/beginner/blitz/cifar10\\_tutorial.html](https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html)



# Implementation of CNN in PyTorch

```
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
print(net)
```

```
Net(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
```

[https://pytorch.org/tutorials/beginner/blitz/cifar10\\_tutorial.html](https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html)

# Implementation of CNN in PyTorch

```
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
print(net)
```

**Increase # of filters in deeper layers:**

As we move forward in the layers, the patterns get more complex; **hence there are larger combinations of patterns to capture.**

```
Net(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
```

[https://pytorch.org/tutorials/beginner/blitz/cifar10\\_tutorial.html](https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html)

# Implementation of CNN in Keras

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

# Number of Parameters

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
<hr/>		
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
<hr/>		
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
<hr/>		
max_pooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
<hr/>		
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
<hr/>		
max_pooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
<hr/>		
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
<hr/>		
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
<hr/>		
flatten_1 (Flatten)	(None, 6272)	0
<hr/>		
dense_1 (Dense)	(None, 512)	3211776
<hr/>		
dense_2 (Dense)	(None, 1)	513
<hr/>		
=====		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		
<hr/>		

Number of parameters in a CONV layer would be :  $((m * n * d)+1)* k$   
**m = shape of filter width**  
**n = shape of filter height**  
**d = number of filters/channels in the previous layer**  
**1 = bias term**  
**k = number of filters**


# CNN

- Video:
  - <https://www.youtube.com/watch?v=3JQ3hYko51Y&feature=youtu.be>
- Interactive visualization:
  - [https://adamharley.com/nn\\_vis/cnn/3d.html](https://adamharley.com/nn_vis/cnn/3d.html)

# Recurrent Neural Network (RNN)

# Recurrent Neural Networks


- Recurrent Neural Networks are networks with loops in them, allowing information to persist.



Recurrent Neural Networks have loops.

In the above diagram, a chunk of neural network, **A**, looks at some input  $x_t$  and outputs a value  $h_t$ .

A loop allows information to be passed from one step of the network to the next.



An unrolled recurrent neural network.

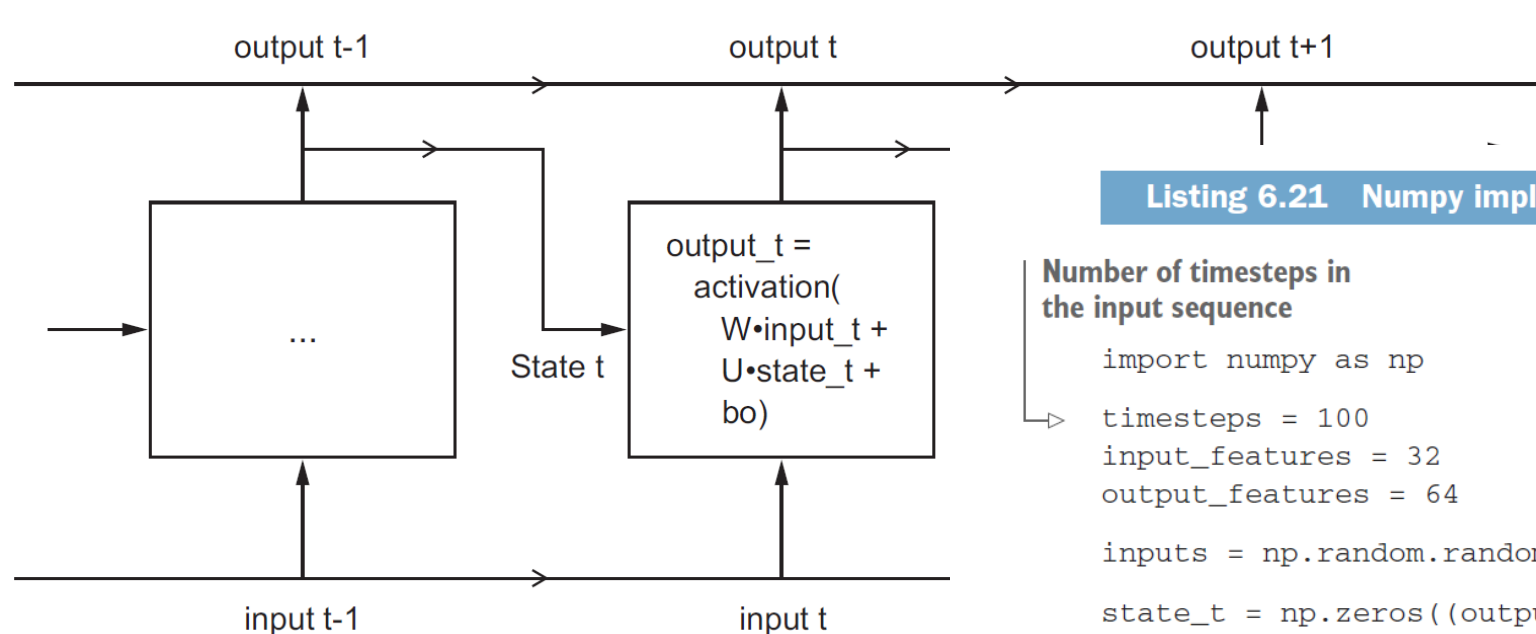
A recurrent neural network can be thought of as multiple copies of the same network, each passing a message to a successor.

The diagram above shows what happens if we unroll the loop.

# Recurrent Neural Networks

- Intuition of Recurrent Neural Networks
  - Human thoughts have persistence; humans don't start their thinking from scratch every second.
    - As you read this sentence, you understand each word based on your understanding of previous words.
  - One of the appeals of RNNs is the idea that they are able to connect previous information to the present task
    - E.g., using previous video frames to inform the understanding of the present frame.
    - E.g., a language model tries to predict the next word based on the previous ones.





**A simple RNN, unrolled over time**

### Listing 6.21 Numpy implementation of a simple RNN

Number of timesteps in the input sequence

```
import numpy as np

timesteps = 100
input_features = 32
output_features = 64
```

Dimensionality of the input feature space

Dimensionality of the output feature space

Input data: random noise for the sake of the example

Initial state: an all-zero vector

```
inputs = np.random.random((timesteps, input_features))

state_t = np.zeros((output_features,))

W = np.random.random((output_features, input_features))
U = np.random.random((output_features, output_features))
b = np.random.random((output_features,))
```

Creates random weight matrices

```
successive_outputs = []
for input_t in inputs:
    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)
```

input\_t is a vector of shape (input\_features,).

```
    successive_outputs.append(output_t)
```

```
    state_t = output_t
```

```
final_output_sequence = np.concatenate(successive_outputs, axis=0)
```

Stores this output in a list

Combines the input with the current state (the previous output) to obtain the current output

The final output is a 2D tensor of shape (timesteps, output\_features).

Updates the state of the network for the next timestep


# RNN/SimpleRNN in PyTorch

```
CLASS torch.nn.RNN(self, input_size, hidden_size, num_layers=1, nonlinearity='tanh', bias=True,  
  batch_first=False, dropout=0.0, bidirectional=False, device=None, dtype=None) [SOURCE]
```



Apply a multi-layer Elman RNN with  $\tanh$  or ReLU non-linearity to an input sequence. For each element in the input sequence, each layer computes the following function:

$$h_t = \tanh(x_t W_{ih}^T + b_{ih} + h_{t-1} W_{hh}^T + b_{hh})$$



```
class RNNTagger(torch.nn.Module):
```

```
    def __init__(self, embedding_dim, hidden_dim, vocab_size, tagset_size):  
        super(RNNTagger, self).__init__()  
        self.hidden_dim = hidden_dim  
  
        self.word_embeddings = torch.nn.Embedding(vocab_size, embedding_dim)  
  
        # The RNN takes word embeddings as inputs, and outputs hidden states  
        # with dimensionality hidden_dim.  
        self.rnn = torch.nn.RNN(embedding_dim, hidden_dim)  
  
        # The linear layer that maps from hidden state space to tag space  
        self.hidden2tag = torch.nn.Linear(hidden_dim, tagset_size)  
  
    def forward(self, sentence):  
        embeds = self.word_embeddings(sentence)  
        rnn_out, _ = self.rnn(embeds.view(len(sentence), 1, -1))  
        tag_space = self.hidden2tag(rnn_out.view(len(sentence), -1))  
        tag_scores = F.log_softmax(tag_space, dim=1)  
        return tag_scores
```

Pytorch reference: <https://pytorch.org/docs/stable/generated/torch.nn.RNN.html>